

BitNet Rust Implementation for Apple Silicon

Project TODO & Roadmap

Project Overview

Create a pure Rust implementation of BitNet (1.58-bit quantized neural networks) optimized for Apple M1/M2/M3 chips, leveraging Metal Performance Shaders and unified memory architecture.

Phase 1: Foundation & Research (Weeks 1-2)

✅ Research & Analysis

☐ Study Reference Implementations

- ☐ Analyze Microsoft BitNet repository structure and algorithms
- ☐ Review MLX framework architecture and Metal integration
- ☐ Examine mlx-bitnet implementation patterns
- ☐ Document key algorithms and data structures

☐ Apple Silicon Architecture Study

- ☐ Research M1/M2/M3 unified memory architecture
- ☐ Understand Metal Performance Shaders (MPS) capabilities
- ☐ Analyze Neural Engine integration possibilities
- ☐ Study Apple's ML Compute framework

☐ Rust Ecosystem Evaluation

- ☐ Evaluate candle-rs vs tch vs burn for tensor operations
- ☐ Research metal-rs bindings and capabilities
- ☐ Assess mlx-rs crate maturity and features
- ☐ Compare SIMD libraries (wide, packed_simd)

✅ Project Setup

☐ Repository Structure

- ☐ Initialize Cargo workspace
- ☐ Set up CI/CD with GitHub Actions
- ☐ Configure benchmarking with criterion
- ☐ Set up documentation with mdbook

☐ Core Dependencies

- ☐ Add tensor computation crate (candle-core)
- ☐ Integrate Metal bindings (metal-rs)
- ☐ Add MLX Rust bindings (mlx-rs)

- ☐ Set up tokenization (tokenizers)
-

Phase 2: Core Implementation (Weeks 3-6)

✅ Quantization Engine

☐ 1.58-bit Quantization

- ☐ Implement BitNet quantization algorithm
- ☐ Create weight quantization functions (-1, 0, +1)
- ☐ Implement activation quantization
- ☐ Add dequantization for computation

☐ Quantization Utilities

- ☐ Weight packing/unpacking functions
- ☐ Quantization-aware training utilities
- ☐ Calibration dataset handling
- ☐ Quantization error analysis tools

✅ BitLinear Layer

☐ Core BitLinear Implementation

- ☐ Implement BitLinear layer as Module trait
- ☐ Add forward pass computation
- ☐ Implement gradient computation for training
- ☐ Add layer normalization integration

☐ Optimization

- ☐ Vectorized operations using SIMD
- ☐ Memory layout optimization
- ☐ Batch processing optimization
- ☐ Cache-friendly data structures

✅ Model Architecture

☐ BitNet Model Structure

- ☐ Implement transformer architecture with BitLinear
- ☐ Add attention mechanism with quantization
- ☐ Implement feed-forward networks
- ☐ Add positional encoding

☐ Model Configuration

- ☐ Create flexible model configuration system
- ☐ Add model serialization/deserialization

- ☐ Implement model loading from checkpoints
 - ☐ Add model validation utilities
-

Phase 3: Apple Silicon Optimization (Weeks 7-10)

✓ Metal Integration

- ☐ **Metal Compute Shaders**
- ☐ Write Metal shaders for BitLinear operations
- ☐ Implement quantized matrix multiplication kernels
- ☐ Add activation function shaders
- ☐ Create memory-efficient data layouts
- ☐ **Metal Performance Optimization**
- ☐ Optimize threadgroup sizes for M1/M2/M3
- ☐ Implement async compute with command buffers
- ☐ Add memory bandwidth optimization
- ☐ Create GPU/CPU hybrid execution paths

✓ Unified Memory Architecture

- ☐ **Memory Management**
- ☐ Implement zero-copy tensor operations
- ☐ Add unified memory pool management
- ☐ Optimize memory allocation patterns
- ☐ Create memory usage profiling tools
- ☐ **Data Pipeline**
- ☐ Streaming data loading for large models
- ☐ Implement prefetching strategies
- ☐ Add memory-mapped model loading
- ☐ Create efficient batch processing

✓ Apple-Specific Features

- ☐ **Neural Engine Integration**
- ☐ Research ANE capabilities for BitNet
- ☐ Implement ANE fallback paths
- ☐ Add performance comparison tools
- ☐ Create hybrid execution strategies
- ☐ **Performance Monitoring**
- ☐ Add Metal GPU performance counters

- ☐ Implement power consumption monitoring
 - ☐ Create thermal throttling detection
 - ☐ Add performance profiling dashboard
-

Phase 4: Inference Engine (Weeks 11-14)

✓ Inference Pipeline

☐ Core Inference Engine

- ☐ Implement forward pass optimization
- ☐ Add batch inference support
- ☐ Create streaming inference for long sequences
- ☐ Implement KV-cache for transformer models

☐ Generation Features

- ☐ Add text generation with sampling strategies
- ☐ Implement beam search and nucleus sampling
- ☐ Add temperature and top-k/top-p controls
- ☐ Create generation stopping criteria

✓ Model Serving

☐ Runtime Optimization

- ☐ Implement model warming strategies
- ☐ Add dynamic batching
- ☐ Create request queuing system
- ☐ Implement load balancing for multi-core

☐ API Interface

- ☐ Create REST API for inference
 - ☐ Add WebSocket support for streaming
 - ☐ Implement authentication and rate limiting
 - ☐ Create client SDKs
-

Phase 5: Training & Fine-tuning (Weeks 15-18)

✓ Training Infrastructure

☐ Training Loop

- ☐ Implement distributed training setup
- ☐ Add gradient accumulation
- ☐ Create checkpointing system

- ☐ Implement learning rate scheduling
- ☐ **Quantization-Aware Training**
- ☐ Add QAT loss functions
- ☐ Implement straight-through estimators
- ☐ Create quantization noise simulation
- ☐ Add quantization regularization

☒ **Fine-tuning Capabilities**

- ☐ **Parameter-Efficient Fine-tuning**
 - ☐ Implement LoRA for BitNet
 - ☐ Add adapter modules
 - ☐ Create prefix tuning support
 - ☐ Implement prompt tuning
 - ☐ **Dataset Handling**
 - ☐ Add common dataset loaders
 - ☐ Implement data preprocessing pipelines
 - ☐ Create data augmentation strategies
 - ☐ Add validation and testing frameworks
-

Phase 6: Testing & Validation (Weeks 19-20)

☒ **Comprehensive Testing**

- ☐ **Unit Tests**
- ☐ Test all quantization functions
- ☐ Validate BitLinear layer correctness
- ☐ Test Metal kernel implementations
- ☐ Verify model loading/saving
- ☐ **Integration Tests**
- ☐ End-to-end inference testing
- ☐ Multi-device testing (M1/M2/M3)
- ☐ Performance regression tests
- ☐ Memory leak detection

☒ **Benchmarking**

- ☐ **Performance Benchmarks**
- ☐ Compare against reference implementations
- ☐ Benchmark across different Apple chips

- ☐ Measure memory usage patterns
 - ☐ Profile inference latency and throughput
 - ☐ **Accuracy Validation**
 - ☐ Validate against original PyTorch models
 - ☐ Test numerical precision
 - ☐ Compare quantization quality
 - ☐ Validate generation quality
-

Phase 7: Documentation & Release (Weeks 21-22)

☒ **Documentation**

- ☐ **Technical Documentation**
- ☐ API documentation with rustdoc
- ☐ Architecture overview
- ☐ Performance tuning guide
- ☐ Troubleshooting guide
- ☐ **User Guides**
- ☐ Quick start tutorial
- ☐ Model conversion guide
- ☐ Fine-tuning tutorial
- ☐ Deployment guide

☒ **Release Preparation**

- ☐ **Package Management**
 - ☐ Prepare crates.io release
 - ☐ Create installation scripts
 - ☐ Add pre-built binaries
 - ☐ Set up package distribution
 - ☐ **Community**
 - ☐ Create example projects
 - ☐ Add contribution guidelines
 - ☐ Set up issue templates
 - ☐ Create community Discord/forum
-

Technical Specifications

Target Performance Goals

- **Inference Speed:** >100 tokens/second on M2 Pro
- **Memory Usage:** <4GB RAM for 7B parameter model
- **Quantization:** 1.58-bit weights, 8-bit activations
- **Model Size:** <2GB for 7B parameter BitNet model

Supported Features

- **Models:** BitNet 1.58, BitNet b1.58
- **Tasks:** Text generation, completion, chat
- **Chips:** M1, M1 Pro/Max, M2, M2 Pro/Max, M3 series
- **Formats:** Safetensors, GGUF, custom BitNet format

Key Dependencies

toml

```
mlx-rs = "0.25"      # Apple MLX framework bindings
candle-core = "0.8"  # Tensor operations
metal-rs = "0.28"    # Metal GPU programming
tokenizers = "0.15"  # Text tokenization
serde = "1.0"        # Serialization
```

Risk Assessment & Mitigation

Technical Risks

- **MLX Rust Bindings Maturity:** Use candle-metal as fallback
- **Metal Shader Complexity:** Start with compute shaders, optimize iteratively
- **Quantization Accuracy:** Validate against reference implementations
- **Memory Constraints:** Implement streaming and model sharding

Timeline Risks

- **Dependency Issues:** Allocate extra time for toolchain setup
 - **Performance Optimization:** Focus on correctness first, optimize second
 - **Apple Silicon Variations:** Test on multiple chip generations early
-

Success Metrics

Performance Metrics

- ☐ Inference speed matches or exceeds MLX Python implementation
- ☐ Memory usage <50% of full precision model
- ☐ Model accuracy within 2% of original BitNet
- ☐ Cold start time <5 seconds for 7B model

Quality Metrics

- ☐ 95%+ test coverage
 - ☐ Zero memory leaks in continuous operation
 - ☐ Comprehensive benchmarking suite
 - ☐ Production-ready documentation
-

Post-Release Roadmap

Short-term (Months 1-3)

- ☐ Community feedback integration
- ☐ Performance optimization based on real-world usage
- ☐ Additional model architecture support
- ☐ Integration with popular inference frameworks

Medium-term (Months 4-6)

- ☐ Multi-modal BitNet support
- ☐ Advanced quantization techniques
- ☐ Edge deployment optimizations
- ☐ Training acceleration features

Long-term (Months 7-12)

- ☐ Research integration with Apple Neural Engine
 - ☐ Advanced model compression techniques
 - ☐ Distributed inference capabilities
 - ☐ Commercial deployment features
-

This roadmap is a living document and will be updated based on progress, feedback, and new developments in the BitNet and Apple Silicon ecosystems.